

# Stat. 654 Homework 1

*Jiayi Liu*

*March 18, 2019*

## Homework 1:

(due Monday March 25, 2019) Read: Chapter 1 and Chapter 2 Problems: Install R and RStudio, if you do not have them installed. Install the R packages `kera` and `tensorflow` for use with a CPU. Run the `cars` Example. Run the `concrete` Example. Run the `iris` Example. Run the code in Chapter 2. A first look at a neural network

## Run the `cars` Example.

### Example: Compare Simple Linear Regression to a single layer NN.

The `cars` dataset in R contains two variables stopping *speed* of cars in mph and *dist* in feet. Using speed to predict stopping distance, two models are fit. See the R code.

- What function is used to normalize the data?
- What percentage of the data is used for *training*? What percentage of the data is used for *testing*?
- What is the fitted linear regression model?
- What is the correlation between the linear regression predicted values and the values from the test data?
- Sketch the NN model that is used to model stopping distance.
- What kind of activation function was used in the ANN? Sketch a picture of what the activation function looks like.
- What is the correlation between the ANN predicted values and the values from the test data?
- Examine the scatterplot of speed by distance with the fitted models. Is the NN fitting a near linear function?
- Which model would you use for prediction? Explain.

### Answer:

Read in data and examine structure.

```
suppressMessages(library("tidyverse"))

## Warning: package 'tibble' was built under R version 3.5.3
## Warning: package 'purrr' was built under R version 3.5.3
## Warning: package 'dplyr' was built under R version 3.5.3
## Warning: package 'forcats' was built under R version 3.5.2
cars <- as.tibble(cars)

## Warning: `as.tibble()` is deprecated, use `as_tibble()` (but mind the new semantics).
## This warning is displayed once per session.
cars


## # A tibble: 50 x 2
##   speed  dist
```

```
##      <dbl> <dbl>
## 1      4      2
## 2      4     10
## 3      7      4
## 4      7     22
## 5      8     16
## 6      9     10
## 7     10     18
## 8     10     26
## 9     10     34
## 10     11     17
## # ... with 40 more rows
```

```
str(cars)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame': 50 obs. of 2 variables:
## $ speed: num 4 4 7 7 8 9 10 10 10 11 ...
## $ dist : num 2 10 4 22 16 10 18 26 34 17 ...
```

```
cars %>% ggplot(aes(x=speed, y=dist)) +
  geom_point(size = 4) +
  ggtitle("Cars data")
```



Apply scaling to entire data frame.

```
cars_norm <- cars %>% mutate(speed = scale(speed), dist=scale(dist))
cars_norm
```

```
## # A tibble: 50 x 2
```


```
##      speed[,1] dist[,1]
##      <dbl>      <dbl>
##  1    -2.16    -1.59
##  2    -2.16    -1.28
##  3    -1.59    -1.51
##  4    -1.59    -0.814
##  5    -1.40    -1.05
##  6    -1.21    -1.28
##  7    -1.02    -0.969
##  8    -1.02    -0.659
##  9    -1.02    -0.348
## 10    -0.832   -1.01
## # ... with 40 more rows
```

```
str(cars_norm)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   50 obs. of  2 variables:
## $ speed: num [1:50, 1] -2.16 -2.16 -1.59 -1.59 -1.4 ...
## ..- attr(*, "scaled:center")= num 15.4
## ..- attr(*, "scaled:scale")= num 5.29
## $ dist : num [1:50, 1] -1.59 -1.28 -1.513 -0.814 -1.047 ...
## ..- attr(*, "scaled:center")= num 43
## ..- attr(*, "scaled:scale")= num 25.8
```

```
cars_norm %>% ggplot(aes(x=speed, y=dist)) +
  geom_point(size = 4) +
  ggtitle("Scaled cars data") +
  scale_x_continuous(limits = c(-2.2, 2)) +
  scale_y_continuous(limits = c(-2, 3))
```

Scaled cars data



Create training and test data.

**Side note:** This is not done using best practices, the `scale()` function should only be applied to the training data not the entire dataset. This is a common practice in many machine learning books. This should be corrected.

```
set.seed(12345)
```

```
idx <- sample(1:50, 40)
```

```
cars_train <- cars_norm[idx, ]
str(cars_train)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':  40 obs. of  2 variables:
## $ speed: num [1:40, 1] 0.681 0.87 1.816 0.87 -0.265 ...
## $ dist : num [1:40, 1] 0.117 0.816 1.631 0.505 -0.271 ...
```

```
cars_test <- cars_norm[-idx, ]
str(cars_test)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':  10 obs. of  2 variables:
## $ speed: num [1:10, 1] -1.589 -1.021 -1.021 -0.454 -0.454 ...
## $ dist : num [1:10, 1] -1.513 -0.969 -0.348 -0.348 0.117 ...
```

```
cars_train <- cars_norm[idx, ]
str(cars_train)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':  40 obs. of  2 variables:
## $ speed: num [1:40, 1] 0.681 0.87 1.816 0.87 -0.265 ...
```

```
## $ dist : num [1:40, 1] 0.117 0.816 1.631 0.505 -0.271 ...
```

```
cars_test <- cars_norm[-idx, ]  
str(cars_test)
```


```
## Classes 'tbl_df', 'tbl' and 'data.frame': 10 obs. of 2 variables:
```

```
## $ speed: num [1:10, 1] -1.589 -1.021 -1.021 -0.454 -0.454 ...
```


```
## $ dist : num [1:10, 1] -1.513 -0.969 -0.348 -0.348 0.117 ...
```

```
cars_train %>% ggplot(aes(x=speed, y=dist)) +  
  geom_point(size = 4) +  
  ggtitle("Training Data") +  
  scale_x_continuous(limits = c(-2.2, 2)) +  
  scale_y_continuous(limits = c(-2, 3))
```

Training Data



```
cars_test %>% ggplot(aes(x=speed, y=dist)) +  
  geom_point(size = 4) +  
  ggtitle("Test Data") +  
  scale_x_continuous(limits = c(-2.2, 2)) +  
  scale_y_continuous(limits = c(-2, 3))
```



Fit a simple linear regression. Train a linear regression model. Predict the Test Data. Compare predicted values with the holdout values.

```
cars_lm <- cars_train %>% lm(dist ~ speed, data = .)
```

```
summary(cars_lm)
```

```
##
## Call:
## lm(formula = dist ~ speed, data = .)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.03373 -0.36619 -0.06137  0.23815  1.66240
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.03134    0.08919  -0.351   0.727
## speed        0.73450    0.09457   7.767 2.31e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5639 on 38 degrees of freedom
## Multiple R-squared:  0.6135, Adjusted R-squared:  0.6033
## F-statistic: 60.32 on 1 and 38 DF, p-value: 2.315e-09
```

```
predicted_lm_dist <- predict(cars_lm, cars_test)

# examine the correlation between predicted and actual values
cor(predicted_lm_dist, cars_test$dist)
```

```
##           [,1]
## [1,] 0.8696118
```

Fit a NN. Train a neural network model. Compare the R code. It is very similar.

```
library(neuralnet)
```

```
## Warning: package 'neuralnet' was built under R version 3.5.3
```

```
##
## Attaching package: 'neuralnet'
## The following object is masked from 'package:dplyr':
##
##      compute
```

```
set.seed(12345)
```

```
cars_model <- cars_train %>% neuralnet(formula = dist ~ speed,
  act.fct = "logistic", hidden = 3, linear.output=TRUE)
```


```
plot(cars_model)
```

Nice plot with the plotnet() function.

```
library(NeuralNetTools)
```

```
## Warning: package 'NeuralNetTools' was built under R version 3.5.3
```

```
par(mar = numeric(4), family = 'serif')
plotnet(cars_model, alpha = 0.6)
```



Predict the Test Data. Compare predicted values with the holdout values.

```
model_results <- compute(cars_model, cars_test[1])
```

```
predicted_dist <- model_results$net.result
```


```
# examine the correlation between predicted and actual values  
cor(predicted_dist, cars_test$dist)
```

```
##           [,1]  
## [1,] 0.8744087
```

Plot the fitted models.

```
ggplot(data=cars_test, aes(x=speed, y=dist)) +  
  geom_point(size = 4) +  
  geom_smooth(method='lm', formula=y~x, fill=NA) +  
  geom_line(aes(y = predicted_dist)) +  
  ggtitle("Test Data Fitted with a Linear Model (blue) and NN (black)") +  
  scale_x_continuous(limits = c(-2.2, 2)) +  
  scale_y_continuous(limits = c(-2, 3))
```





**Example: Compare Simple Linear Regression to a Deep Learning, multilayer neural network.**

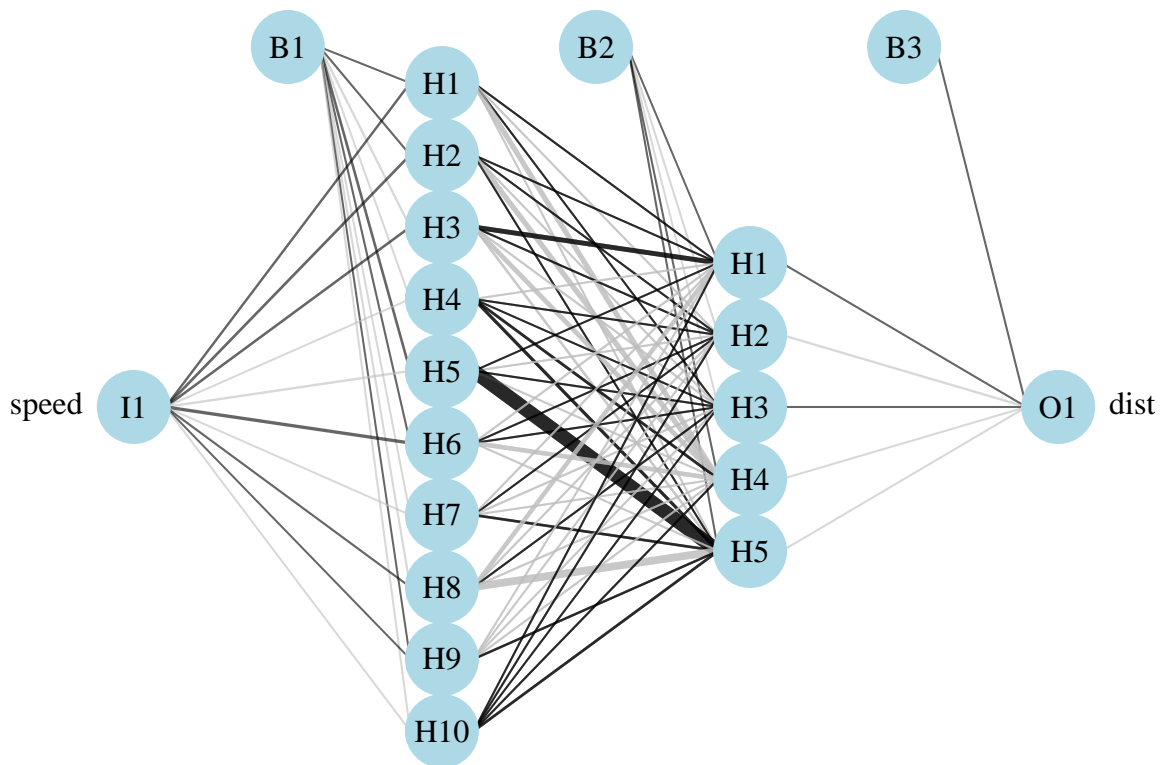
- Do you think this model will overfit?
- What does parsimonious mean?
- Suggest a better measure for goodness-of-fit.

```
cars_model <- cars_train %>% neuralnet(formula = dist ~ speed,
  act.fct = "logistic", hidden = c(10,5), linear.output=TRUE)

plot(cars_model)
```

Nice plot with the plotnet() function.

```
par(mar = numeric(4), family = 'serif')
plotnet(cars_model, alpha = 0.6)
```



Predict the Test Data. Compare predicted values with the holdout values.

```
model_results <- compute(cars_model, cars_test[1])
```

```
predicted_dist <- model_results$net.result
```


```
# examine the correlation between predicted and actual values
cor(predicted_dist, cars_test$dist)
```

```
##           [,1]
## [1,] 0.8890834
```

Plot the fitted models.

```
ggplot(data=cars_test, aes(x=speed, y=dist)) +
  geom_point(size = 4) +
  geom_smooth(method='lm', formula=y~x, fill=NA) +
  geom_line(aes(y = predicted_dist)) +
  ggtitle("Test Data Fitted with a Linear Model (blue) and NN (black)") +
  scale_x_continuous(limits = c(-2.2, 2)) +
  scale_y_continuous(limits = c(-2, 3))
```

Test Data Fitted with a Linear Model (blue) and NN (black)



## Run the concrete Example.

### Step 2: Exploring and preparing the data —

read in data and examine structure

```
concrete <- read.csv("http://www.sci.csueastbay.edu/~esuess/stat654/Poster/concrete.csv")
str(concrete)
```

```
## 'data.frame':  1030 obs. of  9 variables:
## $ cement      : num  141 169 250 266 155 ...
## $ slag        : num  212 42.2 0 114 183.4 ...
## $ ash         : num  0 124.3 95.7 0 0 ...
## $ water       : num  204 158 187 228 193 ...
## $ superplastic: num  0 10.8 5.5 0 9.1 0 0 6.4 0 9 ...
## $ coarseagg   : num  972 1081 957 932 1047 ...
## $ fineagg     : num  748 796 861 670 697 ...
## $ age         : int  28 14 28 28 28 90 7 56 28 28 ...
## $ strength    : num  29.9 23.5 29.2 45.9 18.3 ...
```

custom normalization function

```
normalize <- function(x) {
  return((x - min(x)) / (max(x) - min(x)))
}
```

```
concrete_norm <- as.data.frame(lapply(concrete, normalize))
summary(concrete_norm$strength)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.2664 0.4001 0.4172 0.5457 1.0000
```

```
concrete_train <- concrete_norm[1:773, ]
concrete_test  <- concrete_norm[774:1030, ]
```

### Step 3: Training a model on the data —


```
library(neuralnet)
```

```
# simple ANN with only a single hidden neuron
set.seed(12345) # to guarantee repeatable results
concrete_model <- neuralnet(formula = strength ~ cement + slag +
                             ash + water + superplastic +
                             coarseagg + fineagg + age,
                             data = concrete_train)
```

```
# visualize the network topology
plot(concrete_model)
```

```
# alternative plot
library(NeuralNetTools)
```

```
# plotnet
par(mar = numeric(4), family = 'serif')
plotnet(concrete_model, alpha = 0.6)
```



## Step 4: Evaluating model performance —

```
# obtain model results
model_results <- compute(concrete_model, concrete_test[1:8])
# obtain predicted strength values
predicted_strength <- model_results$net.result
# examine the correlation between predicted and actual values
cor(predicted_strength, concrete_test$strength) # higher than stated in book 0.7170368646
```

```
##           [,1]
## [1,] 0.8064656
```

*# produce actual predictions by*

```
head(predicted_strength)
```

```
##           [,1]
## 774 0.3258992
## 775 0.4677425
## 776 0.2370268
## 777 0.6718811
## 778 0.4663429
## 779 0.4685272
```

```
concrete_train_original_strength <- concrete[1:773,"strength"]
```

```
strength_min <- min(concrete_train_original_strength)
strength_max <- max(concrete_train_original_strength)
```

```
head(concrete_train_original_strength)

## [1] 29.89 23.51 29.22 45.85 18.29 21.86

# custom normalization function
unnormalize <- function(x, min, max) {
  return( (max - min)*x + min )
}

strength_pred <- unnormalize(predicted_strength, strength_min, strength_max)
#strength_pred
```

## Step 5: Improving model performance —

```
# a more complex neural network topology with 5 hidden neurons
set.seed(12345) # to guarantee repeatable results
concrete_model2 <- neuralnet(strength ~ cement + slag +
                             ash + water + superplastic +
                             coarseagg + fineagg + age,
                             data = concrete_train, hidden = 5, act.fct = "logistic")

# plot the network
plot(concrete_model2)

# plotnet
par(mar = numeric(4), family = 'serif')
plotnet(concrete_model2, alpha = 0.6)

# evaluate the results as we did before
model_results2 <- compute(concrete_model2, concrete_test[1:8])
predicted_strength2 <- model_results2$net.result
cor(predicted_strength2, concrete_test$strength) # higher than stated in book 0.801444583

##           [,1]
## [1,] 0.9244533

# try different activation function
# a more complex neural network topology with 5 hidden neurons
set.seed(12345) # to guarantee repeatable results
concrete_model2 <- neuralnet(strength ~ cement + slag +
                             ash + water + superplastic +
                             coarseagg + fineagg + age,
                             data = concrete_train, hidden = 5, act.fct = "tanh")

# evaluate the results as we did before
model_results2 <- compute(concrete_model2, concrete_test[1:8])
predicted_strength2 <- model_results2$net.result
cor(predicted_strength2, concrete_test$strength)

##           [,1]
## [1,] 0.5741729
```

## using h2o deeplearning

```
library(h2o)

## Warning: package 'h2o' was built under R version 3.5.3
##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321
## For more information visit http://docs.h2o.ai
##
## -----
##
## Attaching package: 'h2o'
##
## The following objects are masked from 'package:stats':
##
##   cor, sd, var
##
## The following objects are masked from 'package:base':
##
##   %*%, %in%, &&, ||, apply, as.factor, as.numeric, colnames,
##   colnames<-, ifelse, is.character, is.factor, is.numeric, log,
##   log10, log1p, log2, round, signif, trunc
h2o.init(nthreads=8, max_mem_size="2G")

## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      1 hours 54 minutes
##   H2O cluster timezone:    America/Los_Angeles
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.22.1.1
##   H2O cluster version age:  2 months and 19 days
##   H2O cluster name:        H2O_started_from_R_jiayi_zqn968
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 1.61 GB
##   H2O cluster total cores: 4
##   H2O cluster allowed cores: 4
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:   FALSE
##   H2O API Extensions:      Algos, AutoML, Core V3, Core V4
##   R Version:                R version 3.5.1 (2018-07-02)
```

```

h2o.removeAll() ## clean slate - just in case the cluster was already running

## [1] 0
h2o.init()

## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      1 hours 54 minutes
##   H2O cluster timezone:    America/Los_Angeles
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.22.1.1
##   H2O cluster version age:  2 months and 19 days
##   H2O cluster name:        H2O_started_from_R_jiayi_zqn968
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 1.61 GB
##   H2O cluster total cores: 4
##   H2O cluster allowed cores: 4
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:   FALSE
##   H2O API Extensions:      Algos, AutoML, Core V3, Core V4
##   R Version:                R version 3.5.1 (2018-07-02)

concrete.hex <- h2o.importFile("http://www.sci.csueastbay.edu/~esuess/stat654/Poster/concrete.csv")

##
|
|
|
|=====| 100%

summary(concrete.hex)

## Warning in summary.H2OFrame(concrete.hex): Approximated quantiles
## computed! If you are interested in exact quantiles, please pass the
## `exact_quantiles=TRUE` parameter.

## cement      slag      ash      water
## Min.   :102.0  Min.   : 0.00  Min.   : 0.00  Min.   :121.8
## 1st Qu.:192.1  1st Qu.: 0.00  1st Qu.: 0.00  1st Qu.:164.9
## Median :272.6  Median : 21.92  Median : 0.00  Median :184.9
## Mean   :281.2  Mean   : 73.90  Mean   : 54.19  Mean   :181.6
## 3rd Qu.:349.9  3rd Qu.:142.68  3rd Qu.:118.26  3rd Qu.:191.9
## Max.   :540.0  Max.   :359.40  Max.   :200.10  Max.   :247.0
## superplastic coarseagg  fineagg  age
## Min.   : 0.000  Min.   : 801.0  Min.   :594.0  Min.   : 1.00
## 1st Qu.: 0.000  1st Qu.: 931.7  1st Qu.:730.8  1st Qu.: 7.00
## Median : 6.376  Median : 967.8  Median :779.1  Median : 28.00
## Mean   : 6.205  Mean   : 972.9  Mean   :773.6  Mean   : 45.66
## 3rd Qu.:10.175  3rd Qu.:1029.1  3rd Qu.:824.0  3rd Qu.: 56.00
## Max.   :32.200  Max.   :1145.0  Max.   :992.6  Max.   :365.00
## strength

```



```

## Min.      : 2.33
## 1st Qu.:23.68
## Median :34.40
## Mean    :35.82
## 3rd Qu.:46.10
## Max.     :82.60

splits <- h2o.splitFrame(concrete.hex, 0.75, seed=1234)

dl <- h2o.deeplearning(x=1:8,y="strength",training_frame=splits[[1]],activation = "Tanh",
                      hidden = c(200,200), distribution = "gaussian")

##
|
|
|
|=====
|
|=====
|
|=====
|
|=====| 100%

dl.predict <- h2o.predict(dl, splits[[2]])

##
|
|
|
|=====| 100%

cor(as.vector(dl.predict), as.vector(splits[[2]]$strength))

## [1] 0.9060864

dl@parameters

## $model_id
## [1] "DeepLearning_model_R_1553049245964_3"
##
## $training_frame
## [1] "RTMP_sid_ad96_2"
##
## $activation
## [1] "Tanh"
##
## $seed
## [1] -7.975319e+18
##
## $distribution
## [1] "gaussian"
##
## $x
## [1] "cement"      "slag"        "ash"         "water"
## [5] "superplastic" "coarseagg"   "fineagg"     "age"
##

```

```
## $y
## [1] "strength"
h2o.performance(dl)

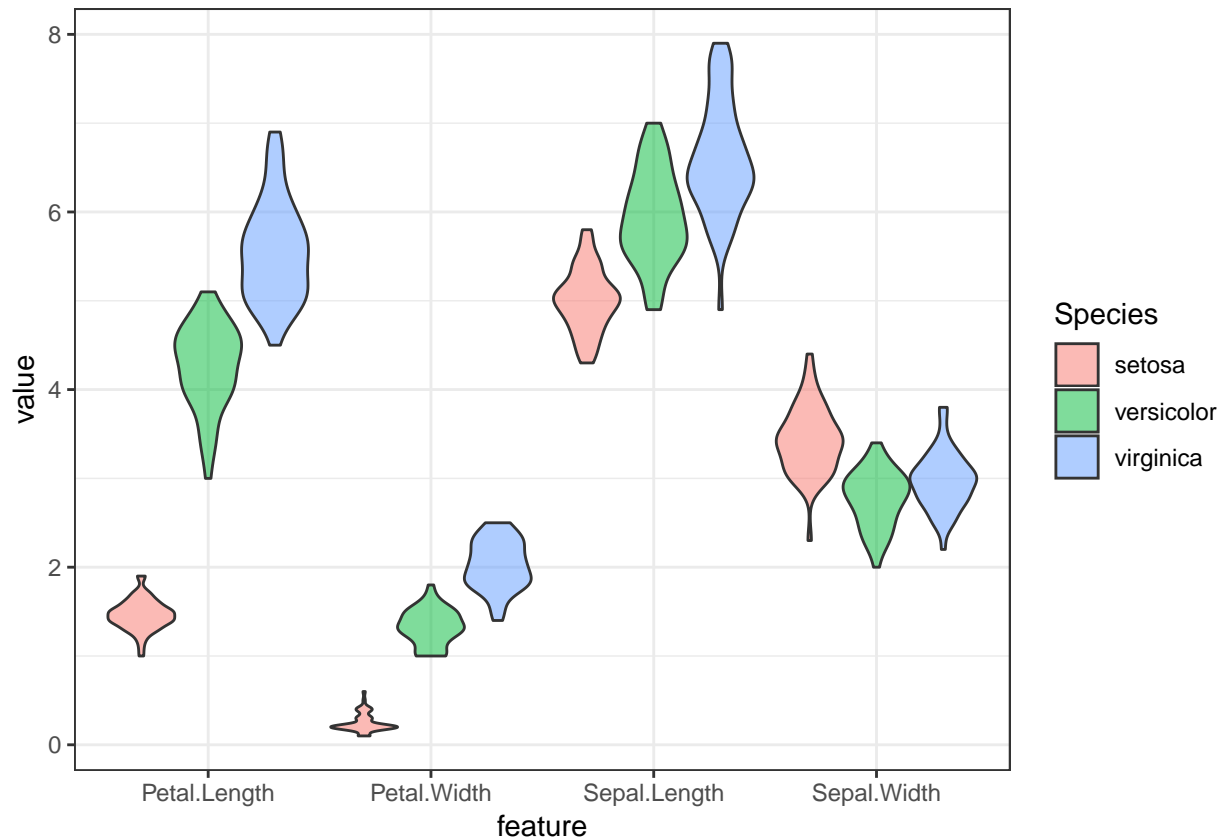
## H2ORegressionMetrics: deeplearning
## ** Reported on training data. **
## ** Metrics reported on full training frame **
##
## MSE: 42.59492
## RMSE: 6.526479
## MAE: 5.150411
## RMSLE: 0.2309562
## Mean Residual Deviance : 42.59492
#h2o.shutdown()
```

## Run the iris Example.

```
library("keras")

## Warning: package 'keras' was built under R version 3.5.3
##
## Attaching package: 'keras'
## The following object is masked _by_ '.GlobalEnv':
##
##      normalize
suppressMessages(library("tidyverse"))

iris %>% as_tibble %>% gather(feature, value, -Species) %>%
  ggplot(aes(x = feature, y = value, fill = Species)) +
  geom_violin(alpha = 0.5, scale = "width") +
  theme_bw()
```



## Prepare data

We start with slightly wrangling the iris data set by renaming and scaling the features and converting character labels to numeric:

```
set.seed(265509)
nn_dat <- iris %>% as_tibble %>%
  mutate(sepal_length = scale(Sepal.Length),
         sepal_width  = scale(Sepal.Width),
         petal_length  = scale(Petal.Length),
         petal_width   = scale(Petal.Width),
         class_label   = as.numeric(Species) - 1) %>%
  select(sepal_length, sepal_width, petal_length, petal_width, class_label)

nn_dat %>% head(3)
```

```
## # A tibble: 3 x 5
##   sepal_length[,1] sepal_width[,1] petal_length[,1] petal_width[,1]
##             <dbl>             <dbl>             <dbl>             <dbl>
## 1          -0.898              1.02            -1.34            -1.31
## 2          -1.14             -0.132            -1.34            -1.31
## 3          -1.38               0.327            -1.39            -1.31
## # ... with 1 more variable: class_label <dbl>
```

Then, we create indices for splitting the iris data into a training and a test data set. We set aside 20% of the data for testing:

```

test_fraction    <- 0.20
n_total_samples <- nrow(nn_dat)
n_train_samples  <- ceiling((1 - test_fraction) * n_total_samples)
train_indices    <- sample(n_total_samples, n_train_samples)
n_test_samples   <- n_total_samples - n_train_samples
test_indices     <- setdiff(seq(1, n_total_samples), train_indices)

```

Based on the indices, we can now create training and test data

```

x_train <- nn_dat %>% select(-class_label) %>% as.matrix %>% .[train_indices,]
y_train <- nn_dat %>% pull(class_label) %>% .[train_indices] %>% to_categorical(3)
x_test  <- nn_dat %>% select(-class_label) %>% as.matrix %>% .[test_indices,]
y_test  <- nn_dat %>% pull(class_label) %>% .[test_indices] %>% to_categorical(3)

```

## Set Architecture

With the data in place, we now set the architecture of our artificial neural network:

```

model <- keras_model_sequential()
model %>%
  layer_dense(units = 4, activation = 'relu', input_shape = 4) %>%
  layer_dense(units = 3, activation = 'softmax')
model %>% summary

```

```

## -----
## Layer (type)                Output Shape          Param #
## -----
## dense_1 (Dense)             (None, 4)             20
## -----
## dense_2 (Dense)             (None, 3)             15
## -----
## Total params: 35
## Trainable params: 35
## Non-trainable params: 0
## -----

```

Next, the architecture set in the model needs to be compiled:

```

model %>% compile(
  loss      = 'categorical_crossentropy',
  optimizer = optimizer_rmsprop(),
  metrics   = c('accuracy')
)

```

## Train the Artificial Neural Network


Lastly we fit the model and save the training progress in the `history` object:

```

history <- model %>% fit(
  x = x_train, y = y_train,
  epochs = 200,
  batch_size = 20,
  validation_split = 0
)
plot(history) +

```

```
ggtitle("Training a neural network based classifier on the iris data set") +
theme_bw()
```



## Evaluate Network Performance

The final performance can be obtained like so:

```
perf <- model %>% evaluate(x_test, y_test)
print(perf)
```

```
## $loss
## [1] 0.2228689
##
## $acc
## [1] 0.9
```

```
classes <- iris %>% as_tibble %>% pull(Species) %>% unique
y_pred <- model %>% predict_classes(x_test)
y_true <- nn_dat %>% pull(class_label) %>% .[test_indices]

tibble(y_true = classes[y_true + 1], y_pred = classes[y_pred + 1],
       Correct = ifelse(y_true == y_pred, "Yes", "No") %>% factor) %>%
  ggplot(aes(x = y_true, y = y_pred, colour = Correct)) +
  geom_jitter() +
  theme_bw() +
  ggtitle(label = "Classification Performance of Artificial Neural Network",
```


```

    subtitle = str_c("Accuracy = ",round(perf$acc,3)*100,"%") +
xlab(label = "True iris class") +
ylab(label = "Predicted iris class")

```

## Classification Performance of Artificial Neural Network

Accuracy = 90%



```
library(gmodels)
```

```
## Warning: package 'gmodels' was built under R version 3.5.3
```

```
CrossTable(y_pred, y_true,
           prop.chisq = FALSE, prop.t = FALSE, prop.r = FALSE,
           dnn = c('predicted', 'actual'))
```

```
##
##
##      Cell Contents
## |-----|
## |                                N |
## |          N / Col Total         |
## |-----|
##
##
##
## Total Observations in Table:  20
##
##
##
##      | actual
## predicted |      0 |          1 |          2 | Row Total |
## -----|-----|-----|-----|-----|
```

```
##          0 |          12 |          0 |          0 |          12 |
##          |          1.000 |          0.000 |          0.000 |          |
## -----|-----|-----|-----|-----|
##          1 |          0 |          3 |          0 |          3 |
##          |          0.000 |          0.600 |          0.000 |          |
## -----|-----|-----|-----|-----|
##          2 |          0 |          2 |          3 |          5 |
##          |          0.000 |          0.400 |          1.000 |          |
## -----|-----|-----|-----|-----|
## Column Total |          12 |          5 |          3 |          20 |
##          |          0.600 |          0.250 |          0.150 |          |
## -----|-----|-----|-----|-----|
##
##
```

## Run the code in Chapter 2. A first look at a neural network

### 2.1 load train and test data

```
library(keras)
mnist <- dataset_mnist()
train_images <- mnist$train$x
train_labels <- mnist$train$y
test_images <- mnist$test$x
test_labels <- mnist$test$y
```

### data prepare

```
str(train_images)

##  int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
str(train_labels)

##  int [1:60000(1d)] 5 0 4 1 9 2 1 3 1 4 ...
str(test_images)

##  int [1:10000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0 ...
str(test_labels)

##  int [1:10000(1d)] 7 2 1 0 4 1 4 9 5 9 ...
train_images <- array_reshape(train_images, c(60000, 28 * 28))
train_images <- train_images/255

test_images <- array_reshape(test_images, c(10000, 28*28))
test_images <- test_images/255

train_labels <- to_categorical(train_labels)
test_labels <- to_categorical(test_labels)
```

## model set

```
network <- keras_model_sequential()%>%
  layer_dense(units=512, activation = "relu", input_shape = c(28*28)) %>%
  layer_dense(units = 10, activation = "softmax")

network %>% compile(
  optimizer= "rmsprop",
  loss = "categorical_crossentropy",
  metrics= c("accuracy")
)
```

## fit model

```
network %>% fit(train_images, train_labels, epochs=5, batch_size=128)

metrics <- network %>% evaluate(test_images, test_labels)
metrics

## $loss
## [1] 0.06781651
##
## $acc
## [1] 0.9798

network %>% predict_classes(test_images[1:10,])

## [1] 7 2 1 0 4 1 4 9 5 9
```

## 2.2 data preparation

```
x <- matrix(rep(0,3*5), nrow = 3, ncol = 5)
x

##      [,1] [,2] [,3] [,4] [,5]
## [1,]    0    0    0    0    0
## [2,]    0    0    0    0    0
## [3,]    0    0    0    0    0

x<- array(rep(0,2*3*2), dim = c(2,3,2))
str(x)

##  num [1:2, 1:3, 1:2] 0 0 0 0 0 0 0 0 0 0 ...

dim(x)

## [1] 2 3 2

length(dim(train_images))

## [1] 2

dim(train_images)

## [1] 60000 784
```



```
typeof(train_images)

## [1] "double"

#digit <- train_images[5,,]
#plot(as.raster(digit, max = 255))
#my_slice <- train_images[10:99,,]
#dim(my_slice)
#my_slice <- train_images[10:99,1:28,1:28]
#dim(my_slice)
#my_slice <- train_images[,15:28,15:28]

#batch <- train_images[1:128,,]
#batch <- train_images[129:256,,]
```

## 2.3 gears of neural networks

```
layer_dense(units=512, activation = "relu")

## <keras.layers.core.Dense>
#output= relu(dot(w, input)+ b)

naive_relu <- function(x){
  for (i in nrow(x))
    for (j in ncol(x))
      x[i,j] <- max(x[i,j],0)
  x
}

naive_relu <- function(x){
  for (i in nrow(x))
    for (j in ncol(x))
      x[i,j] = x[i,j]+y[i,j]
  x
}

x <-array(round(runif(1000,0,9)), dim = c(64,3,32,10))
y <- array(5, dim = c(32,10))
z <- sweep(x, c(3,4), y, pmax)

#z <- x+y
#z <- pmax(z,0)

#sweep(x, 2, y, '+')

naive_vector_dot <- function(x,y){
  z<- 0
  for (i in 1:length(x))
    z <- z+x[[i]]*y[[i]]
  z
}
```

```
naive_matrix_vector_dot <- function(x,y){
  z<- rep(0, nrow(x))
  for (i in 1:nrow(x))
    for (j in 1:ncol(x))
      z <- z[[i]]+x[[i,j]]*y[[j]]
  z
}
```

```
naive_matrix_vector_dot <- function(x, y) {
  z <- rep(0, nrow(x))
  for (i in 1:nrow(x))
    z[[i]] <- naive_vector_dot(x[i,], y)
  z
}
```

```
naive_matrix_dot <- function(x, y) {
  z <- matrix(0, nrow = nrow(x), ncol = ncol(y))
  for (i in 1:nrow(x))
    for (j in 1:ncol(y)) {
      row_x <- x[i,]
      column_y <- y[,j]
      z[i, j] <- naive_vector_dot(row_x, column_y)
    }
  z
}
```

```
train_images <- array_reshape(train_images, c(60000, 28 * 28))
```

```
x <- matrix(c(0, 1,
              2, 3,
              4, 5),
            nrow = 3, ncol = 2, byrow = TRUE)
x <- array_reshape(x, dim = c(6, 1))
x <- array_reshape(x, dim = c(2, 3))
x <- matrix(0, nrow = 300, ncol = 20)
dim(x)
```

```
## [1] 300 20
```

```
x <- t(x)
```

## 2.5 Looking back at our first example

```
network %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy",
  metrics = c("accuracy"))
```

```
compile(
  network,
  optimizer = "rmsprop",
  loss = "categorical_crossentropy",
  metrics = c("accuracy"))
```

```
network %>% fit(train_images, train_labels, epochs = 5, batch_size = 128)
```